

Mastering Data Flow:

Empower Your Projects with Prefect's
Pipeline Magic



Dr Adam Hill

AIM

By the end of this session you will:

- Understand what Prefect is
- Build and execute tasks and flows
- Have scheduled a flow using deployment
- Have a grasp of what else can be done
- Had some fun 😄

PREFECT



Before we get started ...

Download the code

<https://bit.ly/pydata-Prefect-WS>



Run: `docker-compose up --build`

Create free accounts:

- Upstash - <https://upstash.com/>
- MongoDB Atlas - <https://www.mongodb.com/cloud/atlas/register>

Adam Hill

tldr;

**COMPLY
ADVANTAGE**

Senior
Data Scientist

Financial Crime Fighter

Day
Job

Charity

DataKind volunteer

DataDive volunteer; Data
Ambassador; Committee
member



DataKind

Former Royal Society
Entrepreneur in Residence

Recovering astrophysicist

Evangelist

Me

Find me at

@astroadamh

www.linkedin.com/in/adambenhill/

Horsewithapointyhat.com



UNIVERSITY OF
Southampton



THE ROYAL
SOCIETY





Dashboard

Prefect Technologies
data-warehouse

Dashboard

Flow Runs

Flows

Deployments

Work Pools

Blocks

Variables

Automations

Event Feed

Event Webhooks

Artifacts

Settings



Help



Bill Palombi
bill@prefect.io

Flow Runs

1,091 total



11

2

1078

0

0

Flows with failed or crashed runs

orbit-to-bigquery

3h 50m ago

1 ^

orbit-to-bigquery > private-numbat auto-scheduled

Failed

2023/09/01 11:45:07 AM 1m 54s 10 task runs

Deployment @ Orbit to BigQuery

Work Pool @ kubernetes-prd-data-warehouse Work Queue @ default

run-census-sync

3h 53m ago

1 v

main-orchestrator

3h 59m ago

1 v

delete-cloud1-tenant

19h 23m ago

2 v

cloud2-to-data-warehouse

6 v

All tags



8h

24h

1w

Task Runs

39,658

39,649 Completed 99.98%

9 Failed 0.02%


Events

407,844

5,668 Block

454 Worker

401,722 Other



Active Work Pools

kubernetes-legacy-data-warehouse



150 total

Polled
33s ago

Work Queues
●

Late runs
0 (10s avg.)

Completes
98.67% +0.9

kubernetes-prd-data-warehouse



852 total

Polled
18s ago

Work Queues
● ● ● ● ● ● ● ●

Late runs
0 (11s avg.)

Completes
99.65% +1.9

kubernetes-stg-data-warehouse



76 total

Polled
18s ago

Work Queues
● ● ● ● ● ● ● ●

Late runs
0 (10s avg.)

Completes
93.42% -4.3

Calling a task from a flow

Use the `@task` decorator to designate a function as a task. Calling the task creates a new task run:

```
from prefect import flow, task
```

```
@task
def my_task():
    print("Hello, I'm a task")
```

```
@flow
def my_flow():
    my_task()
```

Let's create our first flow!

Adding detail to the task


```
import datetime
from prefect import flow, task

def generate_task_name():
    date = datetime.datetime.now(datetime.timezone.utc)
    return f"{date:%A}-is-a-lovely-day"

@task(name="My Example Task",
      description="An example task for a tutorial.",
      task_run_name=generate_task_name)
def my_task(name):
    pass

@flow
def my_flow():
    # creates a run with a name like "Thursday-is-a-lovely-day"
    my_task(name="marvin")
```


Logging out of the box

 `log_prints=True`

We could have achieved the exact same outcome by using Prefect's convenient `log_prints` keyword argument in the `flow` decorator:

```
@flow(log_prints=True)
def get_repo_info(repo_name: str = "PrefectHQ/prefect"):
    ...
```

repo_info.py

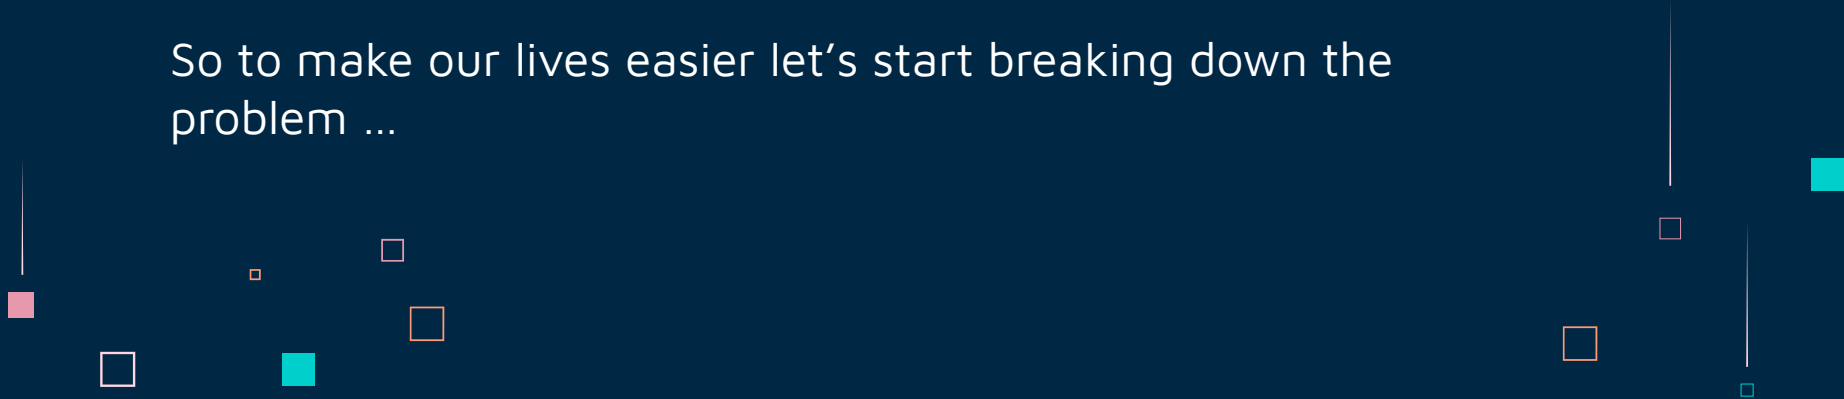
```
import httpx
from prefect import flow, get_run_logger

@flow
def get_repo_info(repo_name: str = "PrefectHQ/prefect"):
    url = f"https://api.github.com/repos/{repo_name}"
    response = httpx.get(url)
    response.raise_for_status()
    repo = response.json()
    logger = get_run_logger()
    logger.info("%s repository statistics 📊:", repo_name)
    logger.info(f"Stars 🌟 : %d", repo["stargazers_count"])
    logger.info(f"Forks 🍴 : %d", repo["forks_count"])
```

Let's start building

Our objective is to build an engine that will monitor regularly for new “tweets” about airlines, move that data into a Kafka message queue. From there we will pick it up, process it, run a sentiment analysis model over it. And finally we will store the original data and the sentiment calculations in a MongoDB in the Cloud!

So to make our lives easier let's start breaking down the problem ...



DEMO / BUILDING




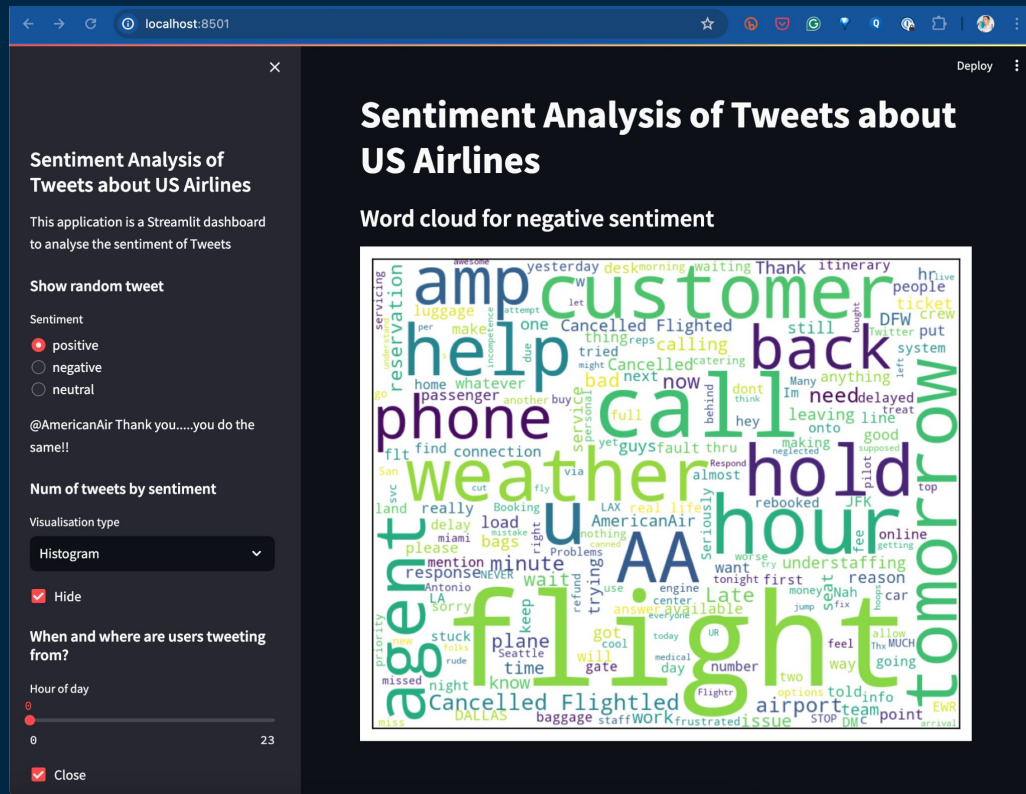
Key problem: Scheduling



```
if __name__ == "__main__":
    get_repo_info.serve(
        name="my-first-deployment",
        cron="* * * * *",
        tags=["testing", "tutorial"],
        description="Given a GitHub repository, logs repository statistics for that
repo.", version="tutorial/deployments",
    )
```

Go to <http://localhost:8501>

Can you adapt the code to pull the data periodically from your new MongoDB that is adding data all the time? 



What we didn't have time for ...

- Parallel processing; e.g. dask can be used to run jobs in parallel
- Secret storage for managing remote deployments
- Plugging into K8s
- Using third-party plug-ins e.g. dbt-runner, docker-runner etc.

Lots more to learn!



We are recruiting @

Comply Advantage



Staff Data Scientist - <https://bit.ly/CA-Staff-DS>



Data Scientist - <https://bit.ly/CA-DataSci>

Adam Hill

tldr;

**COMPLY
ADVANTAGE**

Senior
Data Scientist

Financial Crime Fighter

Thank

you

DataKind volunteer

DataDive volunteer; Data
Ambassador; Committee
member



DataKind

Former Royal Society
Entrepreneur in Residence

Recovering astrophysicist

for

listening

Find me at

@astroadamh

www.linkedin.com/in/adambenhill/

Horsewithapointyhat.com



UNIVERSITY OF
Southampton



THE ROYAL
SOCIETY

